

Convolutional Neural Networks for Image Classification: Analysis and Implementation

Delicia Manasseh a1899076

Abstract—The report is to presents research and to compare various CNN architectures using the CIFAR-10 image classification task. Over the course of the past ten years, numerous CNN applications have been developed for image classification, with a number of diverse proposed architecture variations having emerged such as VGG16, VGG19, and several variants of ResNet. Results have been impressive for deeper models like ResNet-50 on larger datasets, but we want to find the best balance between model complexity and performance under the particular constraints of CIFAR-10. Other exciting aspects dealt with in the project are architectural choices, hyperparameter tuning, and regularization techniques related to CNNs. By systematic experimentation, we arrived at an hyper parameter tuned ResNet-18 featuring a validation accuracy of 85.10%, thus proving that even relatively compact models can provide excellent results when properly optimized. This is important because, in practice, significant computational resources are often lacking. Our findings instead indicate that focused optimization of network design and training strategies can result in far more efficient solutions than pursuing only much deeper architectures. Further, this analysis provides insights into the tradeoffs between model complexity and performance, giving practical guidance towards selecting and tuning CNN architectures on specific project needs.

I. INTRODUCTION

Design choices in the architecture of Convolutional Neural Networks impact performance in image classification. This report investigates how changes in CNN architectural design have influenced performance in image classification over the past decade, a period during which several incrementally amazing architectures have been introduced. In our studies, we used the CIFAR-10 dataset, which contains 50,000 training images and 10,000 test images of size 32x32x3 in 10 classes. While much smaller than databases like ImageNet, that contain upwards of one million training images, CIFAR-10 is none-the-less sufficient in complexity for meaningful architectural comparison. The research focused on six different CNN architectures, each representing a different approach to designing the network. This was accomplished by the MobileNetV2, using depthwise separable convolutions to achieve computational efficiency, the ResNet family, ResNet-18, ResNet-34, ResNet-50 relieving the problem of gradient degradation via skip connections and VGG architectures, VGG-16 and VGG-19, which utilize the more-straightforward approach of stacking convolutional layers with small filters. The purpose of the project is to study various aspects of CNN architecture along with their effect on classification performance. This ranges from an analysis of how the various structural components affect model accuracy, the relation between network depth and

performance, computational demands for different architectures, among others. Such orderly evaluation will offer insights into practical implications of architectural decisions in CNN design. The analysis will revolve around several key aspects:

- It provides an extensive comparison of the known CNN architectures on the CIFAR-10 dataset with respect to their performance capability and computational requirement.
- An improved ResNet-18 design, where the achieved validation accuracy is 85.10%, demonstrating its effectiveness through targeted architectural refinements.
- The empirical investigation shall proceed to describe how various architectures respond to training processes, including a detailed examination of the convergence pattern and accuracy progressions.
- Testing and analysis have been done systematically in order to develop evidence based guidelines on the selection of appropriate CNN structures, based on performance requirements and computational constraints.

This paper is organized as follows: Section 2 details our methodology, including dataset preparation, architecture implementations, and training strategies. Section 3 presents our experimental results and performance comparisons. Section 4 discusses architectural impacts and trade-offs, while Section 5 concludes with recommendations and future research directions.

A. Problem Background

This report currently stands limited to exploring CNN and how model performance can be enhanced in image classification, basically striking a balance between the accuracy-computational efficiency perspective. These CNN architectures have been capable of developing phenomenal capabilities on very large datasets like ImageNet; however, their application to much more constrained situations brings along quite different challenges worth looking at.

- 1) **Architecture Optimization** The adaptation of CNN architectures for the CIFAR10 dataset is unique in its challenges as opposed to the dataset used in standard architecture development. In particular, CIFAR10 contains images of size 32x32x3; this is much smaller than the 227x227x3 format taken by most ImageNet models. This dimensional constraint significantly affects how architectural decisions influence model performance.

While down scaling the input size, standard architectures like VGG, designed for bigger image processing, require careful treatment. Adjustment will be much more than simple rescaling; a deep rethink of the network structure is necessary in order to keep effective feature extraction while avoiding unnecessary high computational overhead. Standard architectures [7] require heavy modification in order to maintain effectiveness in feature extraction across much smaller scales, which is accomplished by reconsidering with thoughtfulness the essential ingredients of the network architecture, such as:

Change in the dimensions of convolutional filters for the input scale Stratus modification of pooling layer placement to preserve critical information While it reduces network complexity to avoid overfitting, discriminative power is preserved.

- 2) **Computational Efficiency** Computational efficiency turns out to be a challenge, especially while considering the optimization of depth in a network. As was underlined by [1], a simple increase in network depth is not a guarantee of better performance; it results in the degradation of both training and validation accuracy. This observation is of crucial importance for such datasets as CIFAR10, when the balance between model complexity and generalization capability becomes a critical factor. Further, [?] stresses that overly complex architecture models will tend to memorize the data rather than learn from features in it-especially on a moderate-scale dataset.
- 3) **Training Optimization** This research aimed to optimize CNN performance on CIFAR-10 [4] by applying dropout regularization [9], batch normalization [2], data augmentation, and adaptive optimizers like Adam [3]. Starting with basic CNN architectures, the researchers enhanced models through these techniques, improving both accuracy and efficiency for practical applications.

Some well-known CNN architectures were tested in this work: VGG16, and VGG19 by Simonyan and Zisserman [7], different ResNet versions such as ResNet-18, ResNet-34, ResNet-50 by He et al., and MobileNetV2 proposed by Sandler et al [8]. After systematic experimentation, the best model turned out to be an improved version of ResNet-18, giving 85.10% on validation with relatively low computational loads (11.7M parameters). This better architecture had dropout layers introduced strategically into the ResNet blocks and thus illustrated that targeted modifications can significantly enhance performance improvement. The experimental methodology involves the inclusion of several key optimization techniques. It uses dropout regularization to prevent overfitting by Srivastava et al. [9], batch normalization to stabilize training by Ioffe and Szegedy [2], and data augmentation to increase model generalization. Adam optimizer by Kingma and Ba [3] was used in training, proving quite effective in traversing the challenging loss landscape of deep architectures. For a model at this scale and complexity, this overall approach to model optimization worked quite well for CIFAR-10. These modifi-

cations are successful and hint at a judicious consideration of both architectural design and training methodology being necessary for optimal CNN performance on such large datasets. The results also introduce targeted enhancements-considering specific dataset characteristics and computational constraints-rather than the mere adaptation of standard architectures.

- **Foundational Architectures:** LeNet-5 established the fundamental CNN structure, while AlexNet [6] demonstrated the potential of deep CNNs, introducing ReLU activations for training efficiency and dropout for regularization.
- **Deep Architectures:** VGG networks [7] utilized small (3x3) filters to emphasize depth but revealed limitations in efficiency and stability, particularly for smaller inputs like CIFAR-10.
- **Residual Learning:** ResNet [1] introduced residual connections, enabling deeper networks to train effectively and influencing subsequent CNN designs. For moderate-scale datasets like CIFAR-10, a balance between depth and complexity is crucial to mitigate overfitting.
- **Efficient Models:** MobileNetV2 [8] achieved computational efficiency through depthwise separable convolutions, making it suitable for resource-constrained environments.
- **Training Optimizations:** Modern CNN training relies on batch normalization [2], dropout [9], and adaptive optimizers like Adam [3] to enhance stability and generalization, especially with limited datasets.
- **Research Gap:** Despite these advancements, adapting CNN architectures for CIFAR-10's unique challenges remains underexplored. This study addresses this gap by evaluating seven CNN architectures, optimizing them for CIFAR-10, and analyzing complexity-accuracy trade-offs.

II. RELATED WORK

The CNN architectures developed to date already show remarkable evolution towards the problem of image classification. However, their optimization for CIFAR-10-like datasets brings out certain aspects on which further research is needed.

A. CNN Architectures Development

Early works in CNN architecture design set the principles still applied in the contemporary models LeCun et al. [5] proposed in LeNet-5 the basic structure of CNN. AlexNet was proposed by Krizhevsky et al. [4] to impose deep learning onto computer vision by means of ReLU activation functions and dropout regularization. The trend toward deeper architectures then gave the VGG-Networks as reviewed by Simonyan and Zisserman [7], based on small-size convolutional filters of size 3×3 employed in a repeating pattern. These networks, though, do suffer from some weaknesses associated with small input dimensions, to which the 32×32 format of CIFAR-10 belongs. A major leap was achieved by ResNet [1], where the obstacles in training very deep networks were resolved by having residual connections, making the process of training much more efficient for deep models. However, many moderate-sized datasets still raise the question concerning the optimal depth.

B. Efficient Architectures and Training Methods

Recent work has revolved around computational efficiency, representative examples including the following: MobileNetV2 by Sandler et al [8]. This model reaches lower computational costs by virtue of depthwise separable convolutions and hence makes the network much more applicable to resource-constrained applications. Current CNN training typically features a few major optimizations such as the following.

- Batch normalization in training by Ioffe and Szegedy [2]
- Dropout regularization by Srivastava et al. to avoid overfitting [9]
- Adaptive optimization by Kingma and Ba for better convergence [3]

C. Current Research Gap

While these advances are indeed very important, the peculiarities of the scale and characteristics of CIFAR-10 remain incompletely addressed with respect to the optimisation of the CNNs. This work attempts to fill this gap by systematically studying architectural choices and their influence on model performance.

III. METHODOLOGY

A. Network Fundamentals and Building Blocks

The fundamental building block of CNNs is the convolution operation. A convolution filter (kernel) slides over the input image, performing element-wise multiplication and summation at each position, as shown in Figure 1.

$$h' = \sum_{ijk} x_{ijk}^r W_{ijk} + b \quad (1)$$

where:

- x_{ijk}^r represents the pixels in the current receptive field
- W_{ijk} are the learnable weights of the filter
- b is the bias term




Fig. 1: Convolution operation: (a) 3x3 filter sliding over input image, (b) Element-wise multiplication and summation process, (c) Resulting feature map generation

B. Dataset

- 1) CIFAR-10 was specifically chosen for this study for several key reasons:
- 2) **Balanced Dataset:** It provides a well-balanced dataset with exactly 6000 images per class, eliminating class imbalance concerns and allowing for fair evaluation across categories.

- 3) **Moderate Scale:** With 60,000 total images, it offers sufficient data for meaningful deep learning experiments while remaining computationally tractable compared to larger datasets like ImageNet.
- 4) **Image Complexity:** The small image size (32x32) presents unique challenges for feature extraction, making it an excellent test case for evaluating CNN architecture modifications.
- 5) **Standardized Benchmark:** As a widely-used benchmark dataset, it enables direct comparison with existing literature and established performance metrics.
- 6) **Real-world Applicability:** The 10 diverse object categories represent common real-world classification tasks, making findings more relevant to practical applications.

C. Dataset and Preprocessing

The CIFAR-10 dataset [10] consists of 60,000 RGB images sized 32x32 pixels, distributed across 10 classes. For our experimentation, we divided the dataset into:

- Training set: 45,000 images
- Validation set: 5,000 images
- Test set: 10,000 images




Fig. 2: Data preprocessing pipeline showing: (a) Original image, (b) Normalized image, (c) Augmented samples

Our preprocessing pipeline includes:

- 1) **Normalization:** Scaling pixel values to [0,1] range by dividing by 255
- 2) **Data Augmentation:**
 - Random rotations (± 15 degrees)
 - Horizontal flips (probability: 0.5)
 - Random width/height shifts ($\pm 10\%$)
 - Random zoom ($\pm 10\%$)

These preprocessing steps are designed to improve model generalization capabilities and mitigate overfitting risks. Batch normalization [2] is implemented throughout the network architecture to stabilize the training process.

D. Model Architectures

The investigation examines several distinct CNN architectures, each representing different approaches to network design: Figure 3 illustrates the key structural differences between these architectures.




Fig. 3: Architecture VGG-16

1) *VGG Architectures*: Both VGG16 and VGG19 follow a straightforward sequential design [7]:

- Sequential arrangement of 3x3 convolutional layers
- Max pooling layers following convolution blocks
- Three fully connected layers
- Progressive filter depth expansion (64, 128, 256, 512)

2) *ResNet Variants*: ResNet architectures (18, 34, and 50 layers) [1] introduce residual connections:

- Basic blocks in ResNet-18/34: Dual 3x3 convolution layers
- Bottleneck blocks in ResNet-50: 1x1, 3x3, 1x1 convolution pattern
- Skip connections implemented every two layers
- Batch normalization following each convolution

3) *MobileNetV2*: Designed for efficiency [8], key features include:

- Depthwise separable convolutions
- Inverted residual blocks
- Linear bottlenecks
- Optimized parameter utilization

TABLE I: Architectural Specifications

Model	Parameters	Depth	FLOPs
VGG16	138M	16	15.5G
VGG19	144M	19	19.6G
ResNet-18	11.7M	18	1.8G
ResNet-34	21.8M	34	3.6G
ResNet-50	25.6M	50	4.1G
MobileNetV2	3.5M	53	0.3G
Improved ResNet-18	11.9M	18	1.8G

IV. EXPERIMENTAL SETUP

A. Training Implementation

The procedure for training was somewhat similar to the methodology described in [He et al.], in that it followed a structured approach to optimization and model evaluation. Each architecture was thus put through systematic training while monitoring carefully the performance metrics and convergence patterns.

B. Training and Validation Curves

For each model, training and validation accuracy and loss curves were generated to track the learning progress and generalization ability over epochs. These plots illustrate how

quickly each model converged and highlight any signs of overfitting.

Figure 10 shows example training and validation accuracy curves for the VGG16, ResNet-50, and Improved ResNet-18 models, demonstrating convergence within 50 epochs. Models with residual connections, such as ResNet-18 and ResNet-50, showed smoother convergence due to skip connections, which stabilize gradient flow. In contrast, VGG models, lacking residual connections, displayed more fluctuation and signs of overfitting despite regularization.

C. Hyperparameter Selection

Hyperparameters were fine-tuned for each model to optimize performance. Key tuning parameters included: **Learning Rate**: Initial learning rates were set to 0.001, with adaptive decay applied to balance early fast learning and later fine-tuning. The learning rate strategy was carefully designed based on empirical observations and theoretical considerations:

- Initial Learning Rate (0.001):
 - Selected based on grid search over 0.01, 0.001, 0.0001
 - 0.01 led to unstable training and divergence
 - 0.0001 resulted in slow convergence
 - 0.001 provided optimal balance between training speed and stability
- Learning Rate Schedule:
 - Implemented reduction on plateau with factor 0.1
 - Patience of 5 epochs to allow for natural fluctuations
 - Minimum learning rate set to 1e-6 to prevent stagnation
- Validation Performance:
 - Tested against constant learning rate
 - Adaptive schedule showed 2.3% improvement in validation accuracy
 - Reduced oscillations in validation loss

Dropout Rate: Dropout rates varied between 0.2 and 0.5 across layers in deeper models like ResNet-50 and Improved ResNet-18, allowing for effective regularization.

- The dropout rates were systematically selected based on the following principles: Progressive Increase: Dropout rates increase with network depth (0.1 in early layers to 0.3 in final layers) because:
 - Earlier layers learn fundamental features that shouldn't be heavily regularized
 - Deeper layers are more prone to overfitting and need stronger regularization
- Empirical Validation: Rates were fine-tuned through experiments:
 - 0.1 dropout after initial convolution blocks: Minimal impact on basic feature extraction
 - 0.2 dropout in middle layers: Balance between regularization and feature preservation
 - 0.3 dropout in final dense layers: Stronger regularization where overfitting typically occurs
- Performance Impact:

- Tested ranges from 0.1 to 0.5 in 0.1 increments
- Rates above 0.3 led to underfitting in early layers
- Rates below 0.1 showed insufficient regularization effect

Batch Size: A batch size of 32 was chosen as the standard across all models, balancing computational efficiency and training stability. These parameters were empirically determined to provide stable training across all architectures while maintaining reasonable computational requirements.

TABLE II: Hyperparameter Tuning and Results for CNN Models on CIFAR-10

Model	Learning Rate Scheduler	Validation
VGG16	Reduction on Plateau (factor 0.1)	61.70%
VGG19	Reduction on Plateau (factor 0.1)	76.19%
ResNet-18	Step Decay (factor 0.1 every 10 epochs)	78.03%
ResNet-34	Step Decay (factor 0.1 every 10 epochs)	77.92%
ResNet-50	Step Decay (factor 0.1 every 10 epochs)	82.94%
MobileNetV2	Reduction on Plateau (factor 0.1)	77.43%
Improved ResNet-18	Reduction on Plateau and Step Decay	85.10%

TABLE III: Comments on CNN Models for CIFAR-10

Model	Comments
VGG16	High dropout rate to prevent overfitting due to large parameter count.
VGG19	Similar to VGG16; deeper layers improve feature extraction, achieving higher accuracy.
ResNet-18	Moderate dropout rate; residual connections improve generalization.
ResNet-34	Similar to ResNet-18, but additional depth did not significantly improve accuracy.
ResNet-50	Deeper architecture and residual connections improve accuracy, minimal overfitting.
MobileNetV2	Lower dropout rate; lightweight architecture achieves efficient training with good accuracy.
Improved ResNet-18	Strategic dropout in residual blocks boosted accuracy; best performance achieved through tuning.

D. Optimization and Learning Rate Scheduling

Optimizer: The Adam optimizer [3] was used because of its adaptive learning rate that allowed not only faster convergence but showed more stability over a number of architectures.

Learning Rate Scheduling:

- Adaptive Reduction
 - Triggered by validation loss plateau
 - Reduction factor: 0.1
 - Applied after three epochs without improvement
 - Helped overcome local minima
- Scheduled Decay
 - Step-wise reduction every 10 epochs
 - Reduction magnitude: factor of 10
 - Provided systematic refinement of parameter updates
 - Facilitated fine-tuning in later training stages.

E. Regularization Techniques

Different complementary regularization techniques were used at the time of implementation with the aim of increasing

model generalization and thereby enhancing the stability of training. Dropout regularization was applied judiciously to network architectures, starting from 0.2 to 0.5, especially within the deeper layers in VGG19 and ResNet-50 models, since overfitting tends to happen in higher-level feature representations. In fact, batch normalization is used across all layers in both ResNet and MobileNetV2 to normalize the activation distribution, increase convergence speed, and allow higher learning rates without training stability degradation. Data augmentation techniques like random rotations, spatial shifts, and horizontal flipping further enhanced the diversity of the training dataset and helped improve the generalization ability of models across different input conditions.

F. Computational Requirements and Training Time

Each model was trained on an [NVIDIA GPU], with training times and computational efficiency noted. Table IV summarizes the average training time per model. Computationally intensive models like VGG16 and VGG19 exhibited longer training times due to their high parameter counts, whereas MobileNetV2 achieved efficient processing, making it suitable for resource-constrained scenarios.

TABLE IV: Average Training Time per Model

Model	Average Training Time (mins)
VGG16	20
VGG19	25
ResNet-18	25
ResNet-34	30
ResNet-50	40
MobileNetV2	22
Improved ResNet-18	30

G. Model Monitoring and Checkpointing

In the meantime, the training process was complemented by systematic monitoring, coupled with model preservation for the best performance. Utilizing the checkpoint mechanism, model states were saved upon improvement of the validation metrics; hence, only the most effective weight configurations were saved during the training process.

Early stopping was utilized to avoid overtraining using a patience threshold of 5 epochs, hence automatically stopping the training procedure if further improvements in validation loss were not detected. This structured model monitoring and preservation of state were effective in ensuring the model is performing optimally without showing overfitting behavior.

V. RESULTS AND ANALYSIS

A. Quantitative Results

The experimental results demonstrate varying performance levels across different architectures. The enhanced ResNet-18 structure achieved superior performance with validation accuracy of 85.10% and test accuracy of 84.20%, notably outperforming other tested architectures.

Confusion Matrix: Figure 12 shows the confusion matrix for Improved ResNet-18, providing insight into class-specific misclassifications.

Training and Validation Curves: Figure 4- 11 presents training and validation accuracy/loss curves, illustrating each model's convergence rate and generalization capability.

B. Comparative Analysis

: Each architecture demonstrated distinct characteristics:

- **VGG Networks**

- VGG16: 61.70% validation accuracy
- VGG19: 76.19% validation accuracy
- Showed greater susceptibility to overfitting
-

- **ResNet Variants**

- ResNet-18: 78.03% validation accuracy
- ResNet-34: 77.92% validation accuracy
- ResNet-50: 82.94% validation accuracy
- Demonstrated consistent performance improvement with depth

- **MobileNetV2**

- Achieved 77.43% validation accuracy
- Balanced efficiency with performance

C. Class-Specific Performance

: The unimproved ResNet-18 demonstrated varying effectiveness across different classes:

- **Strong Performance:**

- Automobile: 94.20
- Frog: 96.10
- Truck: 95.70

- **Challenging Categories:**

- Cat: 60.40
- Dog: 70.50
- Bird: 73.10

These results suggest that model performance correlates strongly with intra-class variation, with more consistent object categories achieving higher accuracy rates.

D. Training Dynamics

The convergence patterns revealed distinct characteristics across architectures:

E. Performance Improvement Analysis

Other critical reasons behind the better performance of the Improved ResNet-18 architecture were a number of main optimization strategies whose relative importance was quantified via ablation studies. Data augmentation was the biggest contributor, where this accounted for 34.9% of the overall gains. This shows the fact that the more diverse the training data, the better the model generalizes. In this respect, the strategic implementation of dropout layers provided the second-largest contribution at 32.5%, showing how targeted regularization is effective in preventing overfitting. The optimization strategy on learning rates added 17.0%, confirming the dual approach of plateau-based and scheduled rate reduction. The remaining 15.6% were the gains contributed by other optimization

techniques, architectural refinements, and adjustments in the training process. It is well illustrated Figure 11 that such a distribution of performance gains provides very good insight into the development of architectures in the future: when models of similar network architectures are dealt with, higher performance can be realized with data-centric approaches and scrupulously implemented regularization techniques.

- ResNet variants showed more stable convergence
- VGG architectures displayed greater fluctuation in validation metrics
- Enhanced ResNet-18 achieved optimal balance between convergence speed and stability

TABLE V: Training and Validation Accuracy for CNN Models on CIFAR-10

Model	Training Accuracy (%)	Validation Accuracy (%)
VGG16	63.00	61.70
VGG19	77.00	76.19
ResNet-18	79.50	78.03
ResNet-34	79.30	77.92
ResNet-50	83.50	82.94
MobileNetV2	78.50	77.43
Improved ResNet-18	86.00	85.10




Fig. 4: Training and Validation Accuracy Curves for VGG16 on CIFAR-10.




Fig. 5: Training and Validation Accuracy Curves for ResNet18 on CIFAR-10.




Fig. 6: Training and Validation Accuracy Curves for ResNet34 on CIFAR-10.




Fig. 7: Training and Validation Accuracy Curves for MobileNetV2 on CIFAR-10.




Fig. 8: Training and Validation Accuracy Curves for CNN Models on CIFAR-10.




Fig. 9: Training and Validation Accuracy Curves for CNN Models on CIFAR-10.




Fig. 10: Training and Validation Accuracy Curves for CNN Models on CIFAR-10.

Distribution of Performance Gains in Improved ResNet-18




Fig. 11: Training and Validation Accuracy Curves for CNN Models on CIFAR-10.




Fig. 12: Confusion Matrix for Improved ResNet-18 on CIFAR-10.

TABLE VI: Validation and Test Accuracy for CNN Models on CIFAR-10

Model	Validation Accuracy	Test Accuracy
Improved ResNet-18	85.10%	84.20%

VI. DISCUSSION AND REFLECTION

A. Design Choices

Throughout this study, various design decisions and trade-offs were made to balance performance and computational efficiency.

- **Key Architectural Decisions:** By the fact that ResNet models used residual connections and MobileNetV2 models chose depthwise separable convolutions, stable gradients and high computational efficiency were achieved.

TABLE VII: Per-Class Accuracy for Improved ResNet-18 on CIFAR-10 Test Set

Class	Correct / Total	Accuracy (%)
Airplane	883/1000	88.30%
Automobile	942/1000	94.20%
Bird	731/1000	73.10%
Cat	604/1000	60.40%
Deer	922/1000	92.20%
Dog	705/1000	70.50%
Frog	961/1000	96.10%
Horse	870/1000	87.00%
Ship	919/1000	91.90%
Truck	957/1000	95.70%

- **Trade-offs Made:** VGG models were deep but at a high computational cost, whereas MobileNetV2 preferred efficiency with losses on accuracy; hence, this paper tunes the improved ResNet-18 to balance between these poles.
- **Challenges Encountered:** Small size of the dataset implied the possibility of overfitting, especially in the VGG models. Introduction of the regularization techniques, dropout, and data augmentations were necessary to deal with those issues.

B. Future Work

Based on this work several directions may be taken:

- **Potential Improvements:** Further optimization over dropout rates, higher batch sizes might allow better generalization, especially in residual networks.
- **Alternative Approaches:** Employing transfer learning from a pretrained ImageNet model could allow for better performance on CIFAR-10, but more on deeper architectures such as VGG19 and ResNet50.
- **Extensions to Other Domains:** These models can be extended for other datasets, such as CIFAR-100 or Tiny ImageNet, in order to test generalizability across similar tasks

VII. CONCLUSION

This work has made an empirical comparison of a number of CNN architectures on CIFAR-10 by demonstrating how architectural choices, hyperparameter tuning, and regularization techniques could affect model performance. Key contributions of this work are as follows:

- **Main Contributions:** for the first time, seven CNN architectures were systematically evaluated, and amongst these, Improved ResNet-18 was found to perform best with 85.10% validation accuracy.
- **Key Findings:** The residual connections and depthwise separable convolutions were regarded as the main factors that achieved good accuracy with sufficiently low computational burden. Dropout techniques and scheduled learning rates helped to avoid overfitting.
- **Broader Implications:** For this reason, this paper recommends the model selection strategy in respect of image classification on a moderate-scale dataset and informs practical usages in diverse areas.

REFERENCES

- [1] He, K., Zhang, X., Ren, S. & Sun, J. 2016, ‘Deep Residual Learning for Image Recognition’, *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 770–778.
- [2] Ioffe, S. & Szegedy, C. 2015, ‘Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift’, *Proceedings of the 32nd International Conference on Machine Learning*, pp. 448–456. Available: <http://proceedings.mlr.press/v37/ioffe15.pdf>
- [3] Kingma, D. P. & Ba, J. 2014, *Adam: A Method for Stochastic Optimization*, arXiv preprint arXiv:1412.6980. Available: <https://arxiv.org/abs/1412.6980>
- [4] Krizhevsky, A. & Hinton, G. 2009, *Learning multiple layers of features from tiny images*, Technical report, University of Toronto. Available: <https://www.cs.toronto.edu/~kriz/learning-features-2009-TR.pdf>
- [5] LeCun, Y., Bottou, L., Bengio, Y. & Haffner, P. 1998, ‘Gradient-Based Learning Applied to Document Recognition’, *Proceedings of the IEEE*, vol. 86, no. 11, pp. 2278–2324. Available: <http://yann.lecun.com/exdb/publis/pdf/lecun-98.pdf>
- [6] Krizhevsky, A., Sutskever, I. & Hinton, G. E. 2012, ‘ImageNet Classification with Deep Convolutional Neural Networks’, *Advances in Neural Information Processing Systems*, pp. 1097–1105. Available: <https://proceedings.neurips.cc/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf>
- [7] Simonyan, K. & Zisserman, A. 2014, *Very Deep Convolutional Networks for Large-Scale Image Recognition*, arXiv preprint arXiv:1409.1556. Available: <https://arxiv.org/abs/1409.1556>
- [8] Sandler, M., Howard, A., Zhu, M., Zhmoginov, A. & Chen, L. C. 2018, ‘MobileNetV2: Inverted Residuals and Linear Bottlenecks’, *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pp. 4510–4520. Available: https://openaccess.thecvf.com/content_cvpr_2018/html/Sandler_MobileNetV2_Inverted_Residuals_CVPR_2018_paper.html
- [9] Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I. & Salakhutdinov, R. 2014, ‘Dropout: A Simple Way to Prevent Neural Networks from Overfitting’, *Journal of Machine Learning Research*, vol. 15, no. 56, pp. 1929–1958. Available: <http://jmlr.org/papers/volume15/srivastava14a/srivastava14a.pdf>
- [10] The University of Toronto, n.d., *CIFAR10 and CIFAR100 datasets*, The University of Toronto, viewed 20 October 2022. Available: <https://www.cs.toronto.edu/~kriz/cifar.html>